

Rich coverage signal and consequences for scaling

Kostya Serebryany
kcc@google.com

July 2023

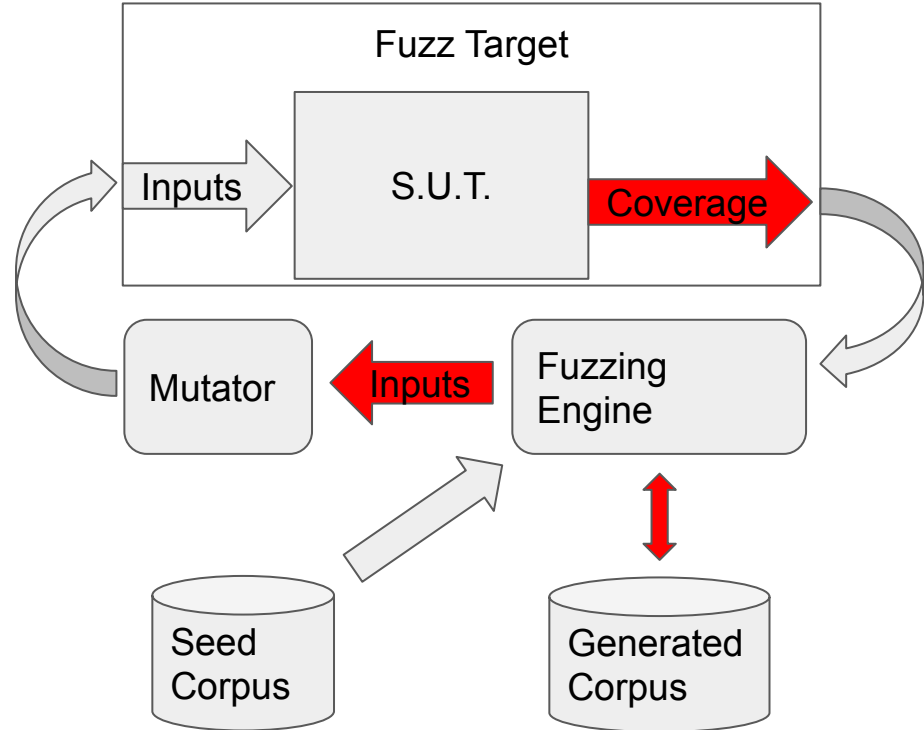
Agenda

- Trivial thoughts on coverage-guided fuzzing
- Case study: SiliFuzz
- Scaling rich coverage signal with the Centipede fuzzing engine

Fuzzing:
generating a maximally diverse
but limited set of test inputs

Coverage-guided fuzzing: coverage is the diversity metric


- Inputs with new **coverage** are added to corpus
- When inputs to mutate are chosen from corpus, **coverage** is taken into account



Guided fuzzing is stronger than unguided (*), but...

Guided fuzzing is unguided most of the time

coverage vs. time




(*) empirically

Initial state: seed corpus



- Input in corpus

Early fuzzing: interesting mutants are added to corpus




• Input in corpus

○ Input not in corpus

● Input newly added to corpus

Corpus grows






• Input in corpus

◦ Input not in corpus

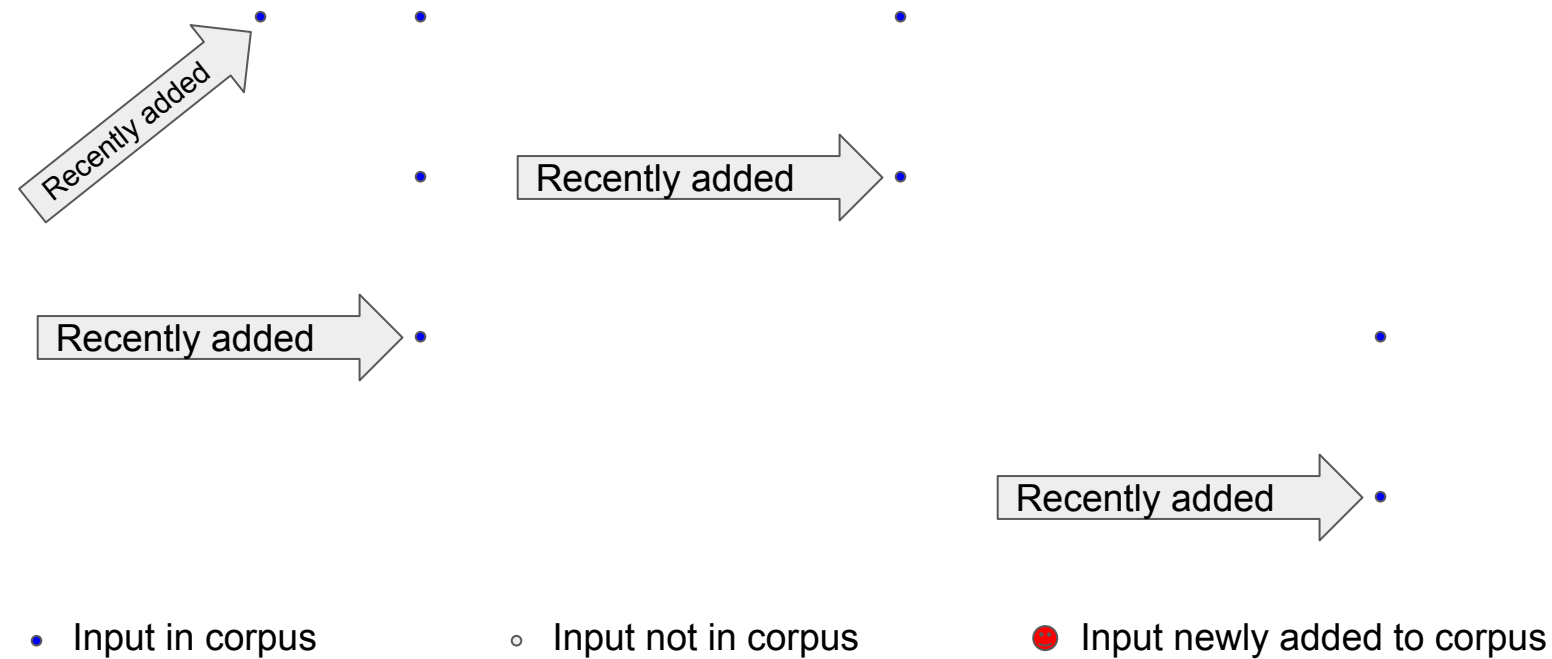
● Input newly added to corpus

More fuzzing




- Input in corpus
- Input not in corpus
- Input newly added to corpus

Corpus grows



And grows ... until it doesn't




• Input in corpus

○ Input not in corpus

● Input newly added to corpus


And grows ... until it doesn't



Problem: traditional coverage signals are sparse

- Most popular coverage signal: control flow edge
 - $O(\text{program size})$ number of different coverage points
 - $O(\text{program size})$ total corpus elements with different edges
- Other popular signals
 - Counters or Value Profiles - still $O(\text{program size})$
 - More numerous than edges and already cause scalability problems

Desired: dense coverage signal, but fuzzing still scales



- Input in corpus

Rich (dense) coverage =>
slower runs, larger corpus, more runs =>
more CPU, RAM, and Disk

Case Study: SiliFuzz

SiliFuzz: detects CPU *bugs* and *defects*

- CPU bug: a family of CPUs is affected
 - [CVE-2021-26339](#): A bug in AMD CPU's core logic may allow for an attacker, using specific code from an unprivileged VM, to trigger a CPU core hang ...
- CPU defect: a single physical CPU is affected, often just one core
 - `imul mem16, %si` leaves `%si` unchanged. Only on one core of one machine.

[0] SiliFuzz: Fuzzing CPUs by proxy. arxiv.org/abs/2110.11519

[1] Cores that don't count. dl.acm.org/doi/10.1145/3458336.3465297 (Google 2021)

[2] Silent Data Corruptions at Scale. arxiv.org/abs/2102.11245 (Facebook 2021)

[3] Detecting silent data corruptions in the wild ([<link>](#), Meta 2022)


SiliFuzz: fuzzing by proxy

- Fuzz something that behaves like a CPU (simulator, RTL design, etc)
 - Generates a diverse set of instruction sequences
 - Slow, requires lots of resources
- Run on one machine per family to detect bugs
 - Separate step (but can be combined with fuzzing too)
 - Relatively fast
- Run on all machines in the fleet to detect defects
 - Separate step
 - Huge scale, very expensive and slow

Anecdotal evidence in support for rich coverage signal

- We fuzzed a CPU proxy for ~ 1 year, found a few bugs & defects, saturated
- Added two “rich signals” and found new bugs within a week
- Added another “rich signal” and found yet another bug
- Rich signals increased corpus sizes by 1000x and more, and caused all sorts of scalability challenges

Centipede: a distributed fuzzing engine with support for rich coverage signal



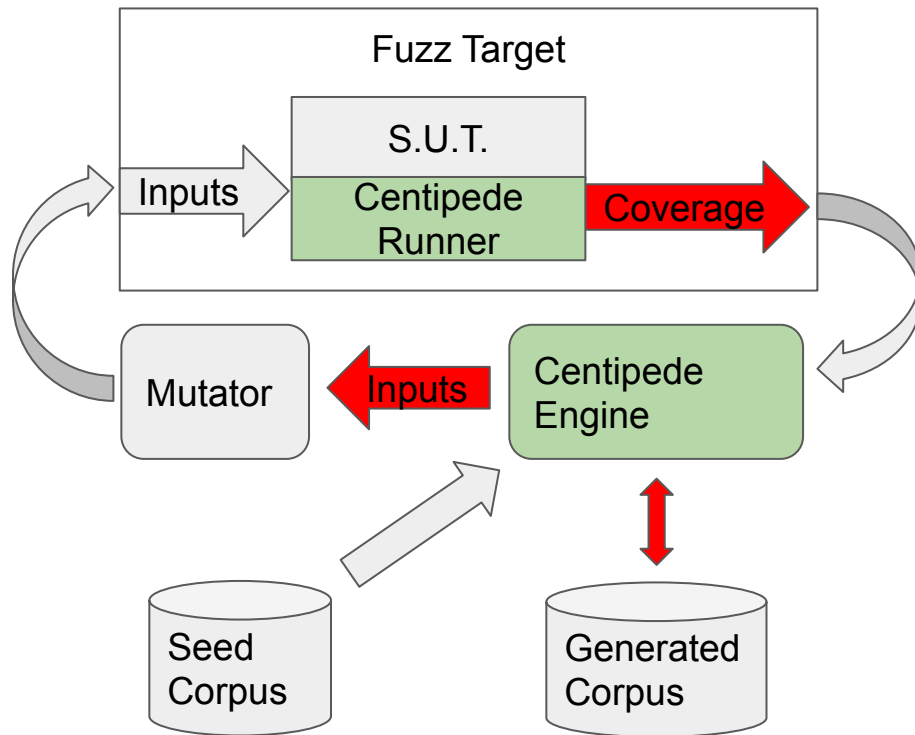
<https://github.com/google/fuzztest/tree/main/centipede>

Centipede's goals

- Efficient fuzzing for any kind and size of target, any language
- libFuzzer-compatible
- In- and out-of-process
- Handle rich 'coverage signal': 1B+ distinct "features"
- Handle corpus of 1B+ inputs
- Massively distributed
- Stateful: restart != recompute
- User can redefine
 - how to execute the target
 - how to mutate inputs

Centipede: Engine vs Runner

- Runner is linked to S.U.T.
 - Consumes a batch of inputs
 - Reports coverage
 - Can be substituted
- Engine is a separate binary
 - Orchestrates mutation & execution
 - Maintains corpus & coverage



Features and Feature Domains

- *Feature* is a 64-bit integer that uniquely represents some program behavior
 - E.g. “this edge has been executed” or “this constant variable has been accessed at this PC”
- A *feature domain* groups all features of the same type
 - E.g. “control flow edges” or “data flow edges”
 - Currently, domain is a fixed integer range of 2^{27}
- For every input, the runner reports a set of features to the engine

Feature domains supported currently

- PCs (control flow edges)
- Call stacks: hash of the top N call stack frames
- Paths: hash of the recent N control flow edges
- PC pairs: a pair of control flow edges
- CMP features: up to 64 different values for every CMP
 - Uses call stacks or paths as context
- Edge counters: 8 buckets, similar to AFL/libFuzzer
- Limited data flows: {global memory location} X {where it is read}
- 16 user-defined feature domains: anything you want
 - [Example from SiliFuzz](#): {instruction opcode} X {index of a bit changed in the register state}
- More will be added: limited only by our imagination
 - Full data flows

Shards, state, distributed execution

- File format: one file contains many data blobs, appendable, remote
- Corpus is represented on disk as N shards
 - file per shard with inputs
 - file per shard with {hash(input), features}
- N workers running concurrently (threads, or different machines)
 - Worker appends only to its own shard, peeks into other shards
- On worker restart, reads inputs and features
 - Recomputes coverage only when unavailable
- Each shard represents only a subset of the corpus

Guided Mutation and Execution Metadata

<not covered here>

Corpus management

- Runner: input => array of features
 - $O(\text{num_inputs} * \text{num_features})$ bytes, too much
- Engine preserves only some of this data
 - Saturated frequencies for all observed features; a feature observed $N+$ times is “frequent”
 - $O(\text{num_features})$ bytes
 - Inputs and their “infrequent” features.
 - $O(\text{num_inputs} * N)$ bytes
- Chooses what inputs to mutate or evict by assigning weights to inputs
 - less frequent feature => better
 - less frequent domain => better
- Future work: apply ML :)

Corpus distillation (minimization)

- Takes N shards (inputs with their features) and produces distilled corpus
 - Same features, minimized number of inputs
- IO-only task, does not involve executing the target

Centipede vs {libFuzzer, AFL, ... }

- With basic coverage: not as tuned, some missing functionality
 - E.g. memcmp / strncmp interceptors
 - But we'll get there
- With rich coverage: will be seen as weaker on short runs (1-2 days)
 - More breadth than depth, requires more iterations to get to the same edge coverage
 - But saturates later
- Works well on huge targets
 - A CPU model: 75M edges (equivalent), 10K+ concurrent shards

Q&A

